

CSI6207

Systems Analysis and Database Design

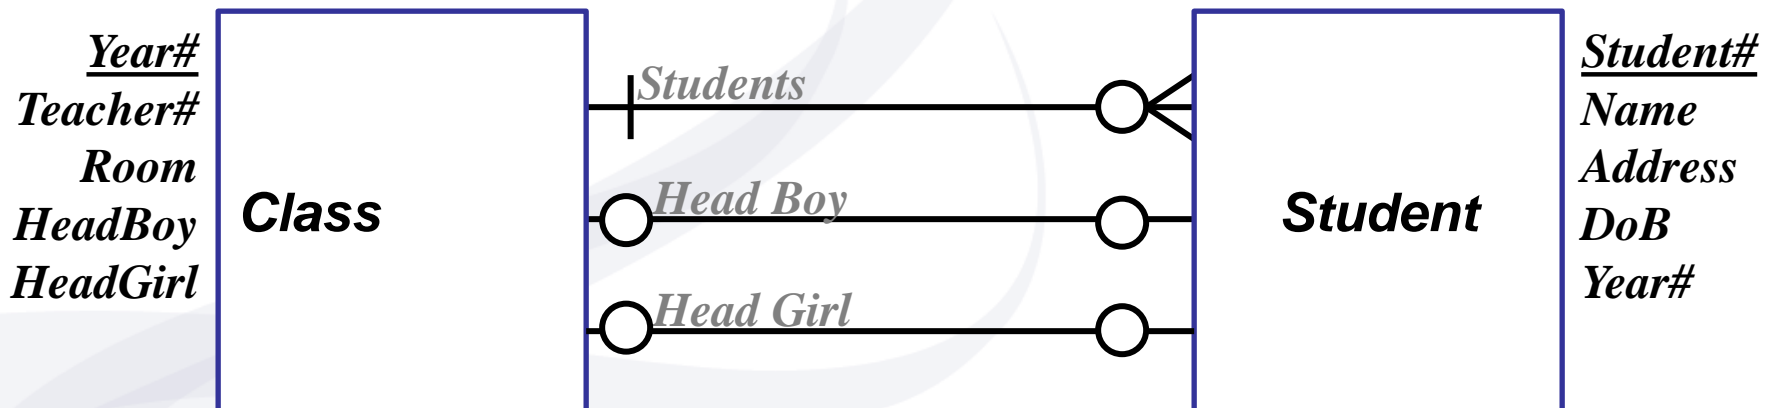
SQL Basics

Objectives

- Entity Relationship (ER) Diagram
 - Multiple and Self-referencing relationships
- Introduction to SQL
- Sample database (used in lectures and labs)
- Simple SELECT statements
- SQL Server & SQL Server Management Studio

- In the examples in the previous lecture, there was only *one relationship between two given entities*
 - This is NOT a rule
- Multiple relationships may occur when *two entities are linked in more than one way*

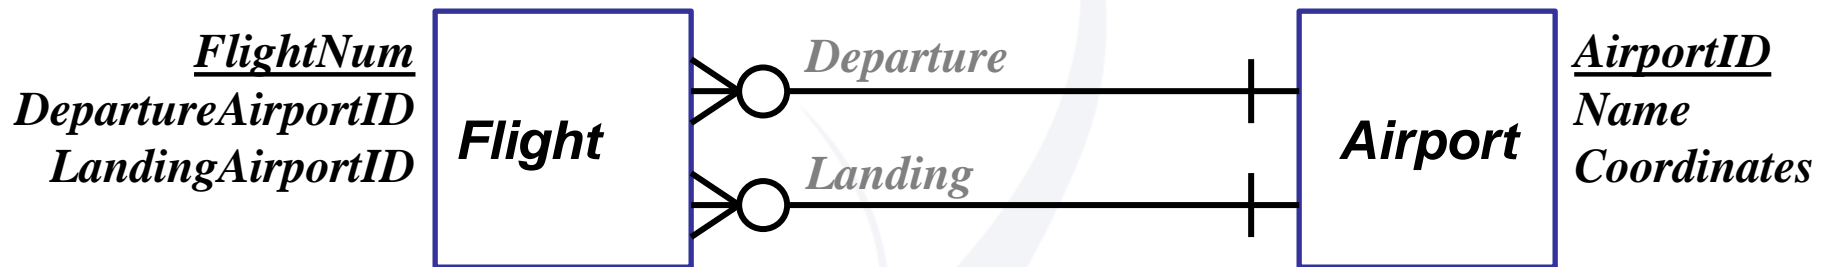
- Imagine a *primary school* database:
 - Each Class (e.g. “Year 6”) has multiple students
 - Each Class has a Head Boy
 - Each Class has a Head Girl
 - Each year *may have* a head boy and a head girl, and *it is not required* for a *student* to be a head boy or girl



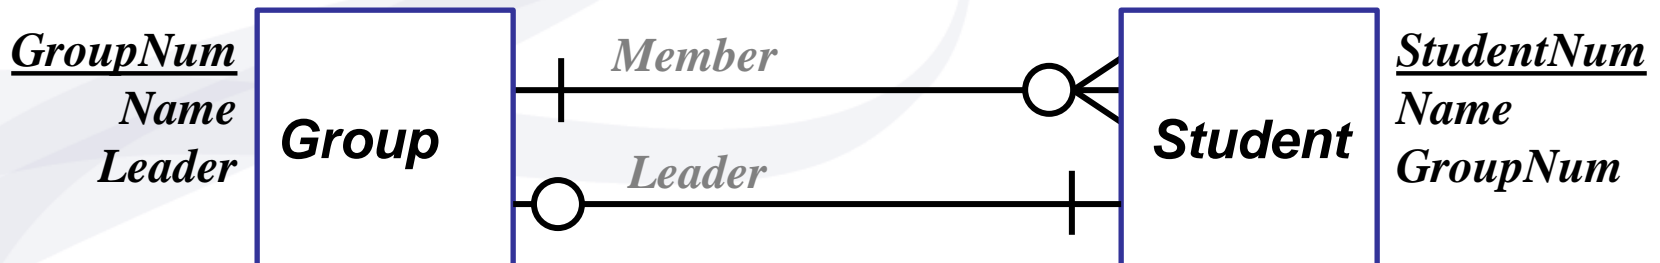
- When multiple relationships exist, it is important to give FKs *meaningful* names, and consider naming the relationships

- Other examples of multiple relationship scenarios:

- “A flight departs from one airport and lands at another”

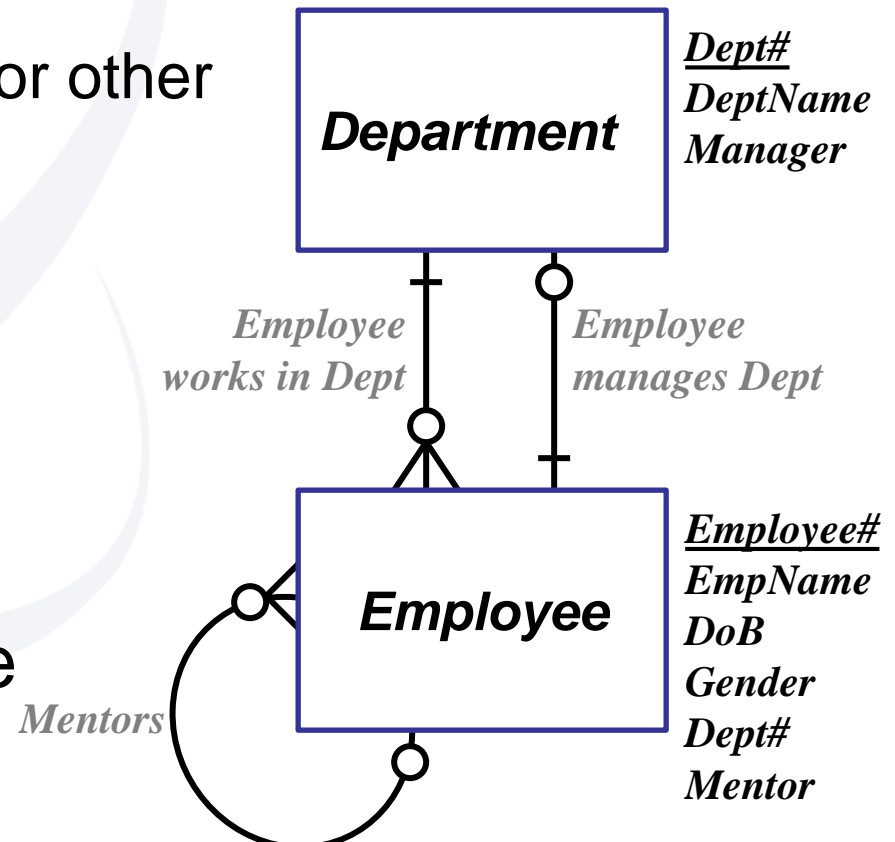


- “Groups can have many students and must have a leader”




Self-Referencing Relationships

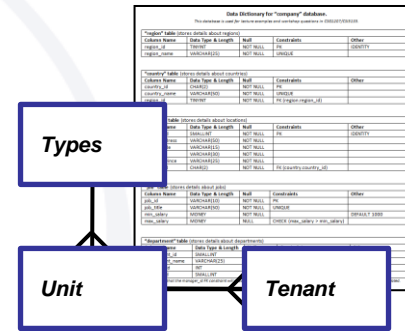
- A company database, with departments and employees
 - One department may have multiple employees
 - Each department also has a department manager
 - Employees may also mentor other employees in a 1:M way
- This is a self-referencing relationship
- Employee has a recursive relationship with itself




From Concept to Database



**Model of system
in client's mind**



**Full database design
(including data dictionary)**



**Physical E-R model
of client's model
(in 3NF or higher)**



**Tables on disk
in a RDBMS**

Relational Database Terminology Recap

department_id	department_name	manager_id	location_id
10	Administration	16	3
20	Marketing	17	4
30	Shipping	7	2
40	IT	4	1
50	Sales	12	5
60	Executive	1	3
70	Accounting	19	3
80	Contracting	NULL	3

Table (*Relation*).

Row (Record).

Column (Attribute).

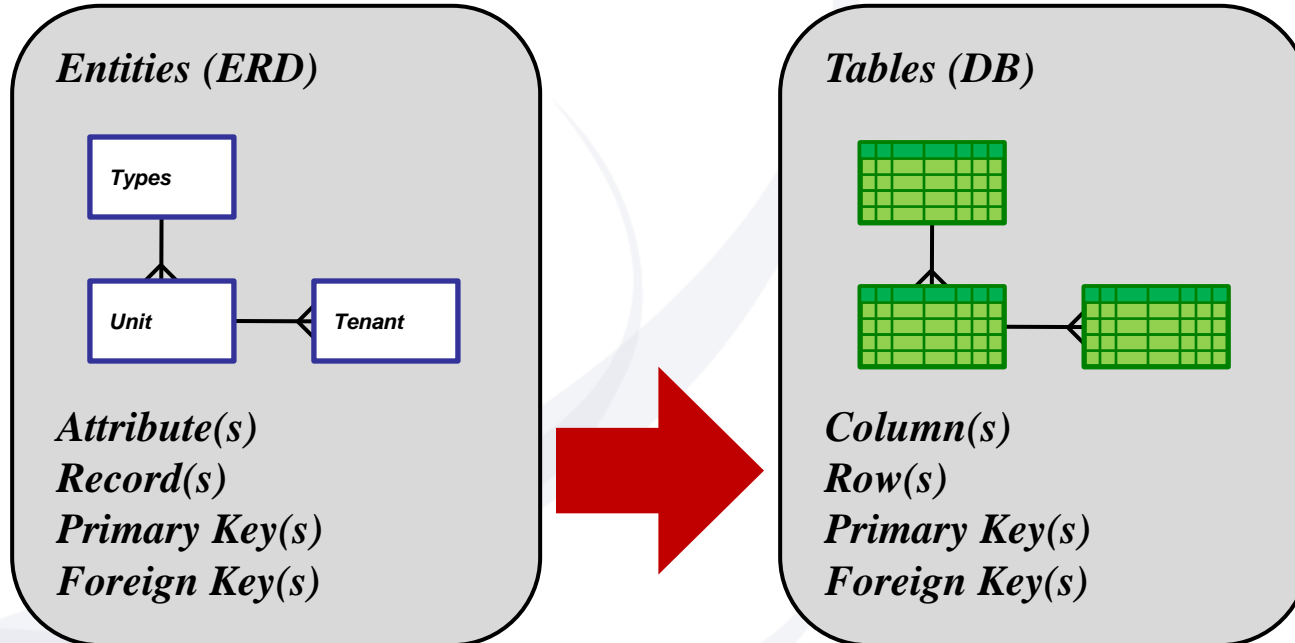
Primary Key.

Foreign Keys.

Field.

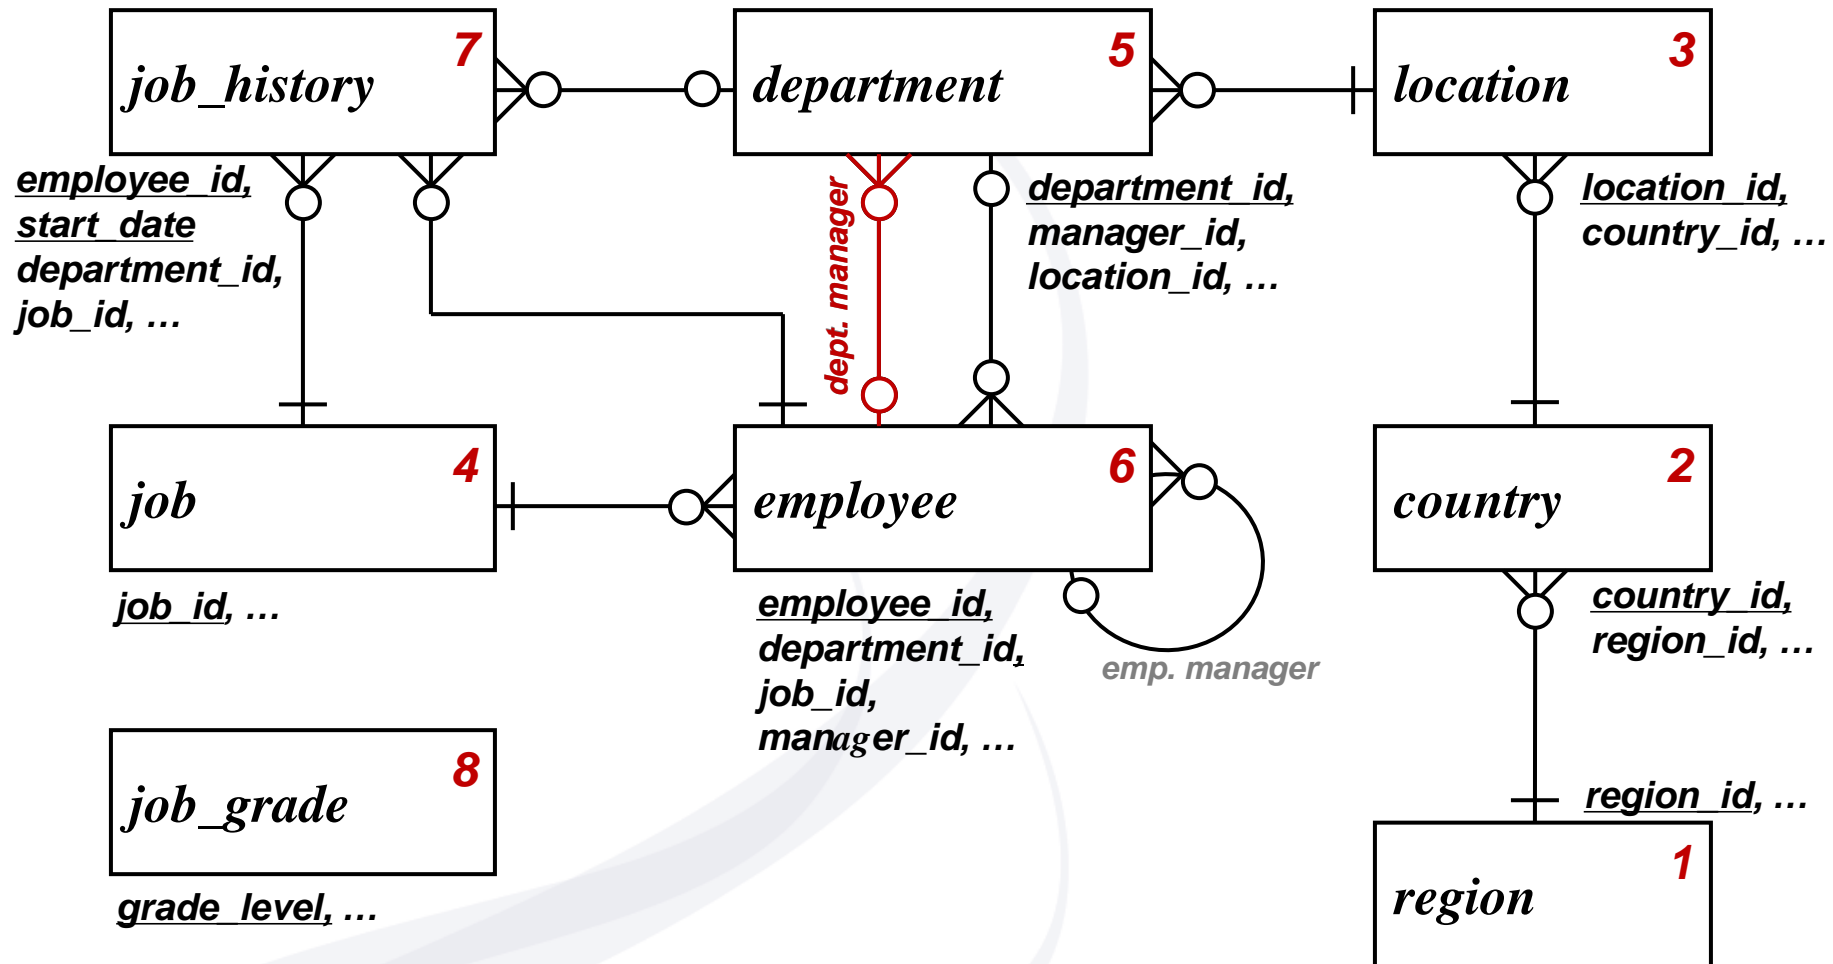
Null (*No Value*).

- Each entity in a physical ERD *translates directly* to a table



- Once a physical model of a normalised system has been created, it's time to *implement it as an actual database*
- The first step is to **create the tables** of your database
- Each entity in your physical model will become a table
- The **order** of **table creation** is important, in order to ensure the existence of **primary** and **foreign** keys in relationships
 - *You cannot have a foreign key column that refers to a table that has not yet been created*

Sample Database – “Company”



- job_grade has no FKs – creation order not important
- The relationship in red is added *after creation and population* to prevent issues with circular references

The Company Database

employee

employee_id	first_name	last_name	gender	email	phone_number	hire_date	job_id	salary	commission_pct	manager_id	department_id
1	Steven	King	M	SKING	515 123 4567	1987-06-17	AD_PRES	24000.00	NULL	NULL	60
2	Neena	Kochhar	F	NKOCHHAR	515 123 4568	1989-09-21	AD_VP	17000.00	NULL	1	60
3	Lex	De Haan	M	LDEHAAN	515 123 4569	1993-01-13	AD_VP	17000.00	NULL	1	60
4	Alexander	Hunold	M	AHUNOLD	590 423 4567	1990-01-03	IT_PROG	9000.00	NULL	3	40
5	Bruce	Ernst	M	BERNST	590 423 4568	1991-03-21	IT_PROG	6000.00	NULL	4	40
6	Diana	Lorentz	F	DLORENTZ	590 423 5567	NULL	IT_PROG	4200.00	NULL	4	40
7	Kevin	Mourgos	NULL	KMOURGOS	650 123 5234	1999-05-07	ST_MAN	6000.00	NULL	1	30
8	Trenna	Rajs	F	TRAJS	650 121 8009	1995-10-17	ST_CLERK	3500.00	NULL	7	30
9	Curtis	Davies	M	CDAVIES	650 121 2994	NULL	ST_CLERK	3100.00	NULL	7	30
10	Randall	Matos	NULL	RMATOS	650 121 2874	1998-03-15	ST_CLERK	2600.00	NULL	7	30
11	Peter	Vargas	M	PVARGAS	650 121 2004	1999-05-07	ST_CLERK	2500.00	NULL	7	30
12	Eleni	Zlotkey	F	EZLOTKEY	(011) 44 1344 429018	2000-01-29	SA_MAN	10500.00	0.20	1	50
13	Ellen	Abel	F	EABEL	(011) 44 1644 429267	1996-05-11	SA_REP	11000.00	0.30	12	50
14	Jonathon	Taylor	M	JTAYLOR	(011) 44 1644 429265	1998-03-24	SA_REP	8000.00	0.20	12	50
15	Kimberely	Grant	F	KGRANT	(011) 44 1644 429263	1999-05-07	SA_REP	7000.00	0.15	12	NULL
16	Jennifer	Whalen	F	JWHALEN	515 123 4444	1987-09-17	AD_ASST	4400.00	NULL	2	10
17	Michael	Hartstein	NULL	MHARTSTE	515 123 5555	1996-02-17	MK_MAN	13000.00	NULL	1	20
18	Pat	Fay	F	PFAY	603 123 6646	1997-08-17	MK_REP	3500.00	NULL	17	20
19	Shelley	Higgins	F	SHIGGINS	515 123 8080	1994-06-07	AC_MGR	12000.00	NULL	2	70
20	William	Gietz	M	WGIEZT	515 123 8181	1994-06-07	AC_ACC...	8000.00	NULL	19	70

The Company Database

department

department_id	department_name	manager_id	location_id
10	Administration	16	3
20	Marketing	17	4
30	Shipping	7	2
40	IT	4	1
50	Sales	12	5
60	Executive	1	3
70	Accounting	19	3
80	Contracting	NULL	3


job

job_id	job_title	min_salary	max_salary
AC_ACCOUNT	Public Accountant	4200.00	9000.00
AC_MGR	Accounting Manager	8200.00	16000.00
AD_ASST	Administration Assistant	3000.00	6000.00
AD PRES	President	20000.00	NULL
AD_VP	Administration Vice President	15000.00	30000.00
IT_PROG	Programmer	4000.00	10000.00
MK_MAN	Marketing Manager	9000.00	15000.00
MK_REP	Marketing Representative	4000.00	9000.00

location

location_id	street_address	postal_code	city	state_province	country_id
1	2014 Jabberwocky Rd	26192	Southlake	Texas	US
2	2011 Interiors Blvd	99236	South San Francisco	California	US
3	2004 Charade Rd	98199	Seattle	Washington	US
4	460 Bloor St. W.	ON M5S 1X8	Toronto	Ontario	CA
5	Magdalen Centre, The Oxford Science Park	OX9 9ZB	Oxford	Oxford	UK

- The basic rule to remember when creating tables is:
“The one side of a one-to-many relationship must always be made before the many side”



- Creating **Unit**, then **Student**, then **Enrolment** would also be appropriate
 - as long as Unit and Student are created before Enrolment

- Deleting a table in a database is known as **dropping** it
- The rule to remember for this is:
 “All tables with foreign keys be dropped before the table they reference is dropped”
- Essentially...
 “The many side of a one-to-many relationship must always be dropped before the one side”
- In case it isn't obvious...
 “The drop order is the reverse of the creation order”

- Once a **well-structured** and **normalised database** has been designed via normalisation and/or ER modelling, it is almost ready to **implement** it as an **actual database** in a DBMS
- The last step in a good design is to create a **data dictionary**
- A data dictionary should contain **all the information** needed to implement the database in a DBMS. This includes:
 - The **names** of all entities and their **attributes**
 - The **domain** of all attributes (**data types**, **constraints**, etc)
 - Details of all **primary** and **foreign** keys
 - Written **descriptions** of entities, attributes, relationships, etc, where needed (e.g. for anything confusing or ambiguous)

- Data dictionaries typically take the form of a number of tables – one table per entity
 - **Order the tables** in an appropriate table creation **order**, or remember to specify this information in the data dictionary

“customer” table (stores details about customers)

Column Name	Data Type & Length	Null	Constraints	Other
customer_id	INT	NOT NULL	PK	IDENTITY
name	VARCHAR(50)	NOT NULL		
phone	VARCHAR(20)	NOT NULL		
address	TEXT	NOT NULL		

“order” table (stores details about customers’ orders)

Column Name	Data Type & Length	Null	Constraints	Other
invoice_id	INT	NOT NULL	PK	IDENTITY
order_date	DATETIME	NOT NULL		
customer_id	INT	NOT NULL	FK (customer.customer_id)	

“item” table (stores details about items for sale)

Column Name	Data Type & Length	Null	Constraints	Other
item_id	INT	NOT NULL	PK	IDENTITY
description	VARCHAR(50)	NOT NULL		
unit_price	MONEY	NOT NULL		

“order_item” table (stores details about the items in an order)

Column Name	Data Type & Length	Null	Constraints	Other
invoice_id	INT	NOT NULL	PK, FK (order.invoice_id)	
item_id	INT	NOT NULL	PK, FK (item.item_id)	
qty	SMALLINT	NOT NULL		DEFAULT 1

- A data dictionary should contain *everything* that someone *needs* to know to *implement* the database in a DBMS
- Some columns of the data dictionary refer to data types and constraints used in the database/DBMS itself
 - We will cover this in upcoming weeks

- Structured Query Language (SQL) is the language used to send **commands** to a database in a **RDBMS**, including...
 - Commands to **retrieve** data from a database
 - (Standard SQL queries using the “**SELECT**” command)
 - Commands to **insert, update or delete data** in a database
 - (Data Manipulation Language - DML)
 - Commands to **create, modify and delete database schemas**
 - (Data Definition Language – DDL)
 - Commands to manage users access control to a **database**
 - (Data Control Language – DCL)
- All these languages (DML, DDL, etc) are part of SQL, and have consistent syntax style and structure
 - They are defined only by their purpose

- SQL is a standardised language supported by just about every RDBMS, but many “variations” exist
 - While the common/basic syntax for most commands remains the same, **some commands have different syntax**
 - They also **add features** which are often only supported by certain products who have implemented that variation
 - It is **unwise** to **rely** heavily on such **features**, as this limits your ability to transfer your database from one DBMS to another

- Here is a list of some common SQL commands...

Data Retrieval	SELECT
Data Manipulation Language (DML)	INSERT UPDATE DELETE
Data Definition Language (DDL)	CREATE DATABASE/TABLE ALTER DATABASE/TABLE DROP DATABASE/TABLE RENAME DATABASE/TABLE TRUNCATE TABLE
Transaction Control	BEGIN/SAVE TRANSACTION COMMIT ROLLBACK
Data Control Language (DCL)	CREATE/ALTER/DROP USER GRANT REVOKE

- A table creation statement in SQL consists of several basic elements:
 - The words **CREATE TABLE**
 - The *name* of the table
 - An *opening parenthesis*
 - *Column definitions* (separated by commas)
 - Includes name, data type and other properties of each column
 - *Constraint definitions* (separated by commas)
 - e.g. Key fields, unique fields...
 - A *closing parenthesis*
 - A SQL terminator (;)

- Remembering that each entity in your physical ERD maps to a table, let's make a "student" table
 - Some DBMSs **do not support** "#" in a field name, so we use alphanumeric characters and underscores only

Student

Student#
FirstName
Surname
Gender
DoB
Phone
Email
Height

CREATE TABLE student

(

student_id	INT	NOT NULL	PRIMARY KEY,
first_name	VARCHAR(20)	NOT NULL,	
surname	VARCHAR(20)	NOT NULL,	
gender	CHAR(1)	NOT NULL,	
dob	DATE	NOT NULL,	
phone	VARCHAR(10)	NULL,	
email	VARCHAR(50)	NULL,	
height	NUMERIC(3,2)	NULL	

);

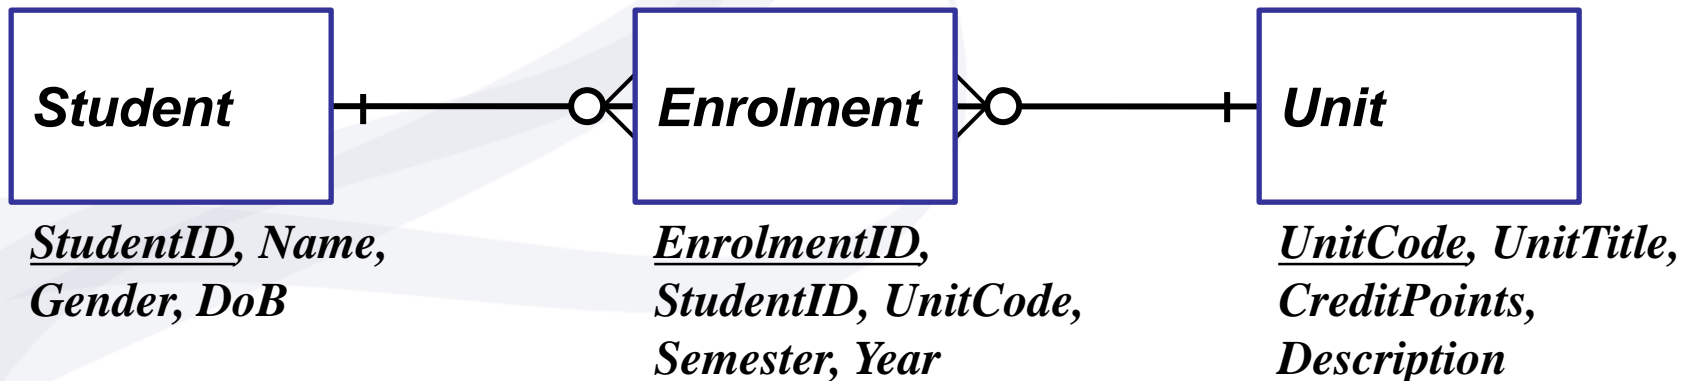
- Since we just discussed it, let's start with something simple: Dropping tables in a database using SQL statements
- The syntax for this statement is simply:
 - **DROP TABLE <table name>;**
- In regards to our last example...
DROP TABLE ApplianceOwner;
DROP TABLE Appliance;
DROP TABLE Tenant;
DROP TABLE Unit;
- Each statement ends with a ; (semicolon)

- Here's another example, this time with a **compound key**...

```
CREATE TABLE order_item  
(  
    invoice_id INT NOT NULL,  
    item_id CHAR(10) NOT NULL,  
    qty TINYINT NOT NULL DEFAULT 1,  
    CONSTRAINT order_item_pk PRIMARY KEY (invoice_id, item_id)  
);
```

- Notes:
 - Both the invoice and **item ids** are part of the **primary key** (we have omitted foreign keys in this example)
 - Item_id is CHAR(10) rather than an INT... Why?
 - Quantity is TINYINT, which accepts anything from 0 to 255, and it has a default value of 1 – used if no quantity specified

- The **foreign key**, or *referential integrity constraint*, designates a column, or combination of columns, as a foreign key and establishes a **relationship** to a **primary key** (or a unique key) in another table (or even the same table).
- **Create order** must be followed – primary key must already exist in order to create foreign key constraint with it
- Example - create tables for the simple enrolment model:



- Student table and Unit table must be created first

```
CREATE TABLE student
```

```
(  
    student_id      INT          NOT NULL    PRIMARY KEY,  
    name           VARCHAR(50)  NOT NULL,  
    gender         CHAR(1)     NOT NULL,  
    dob            DATE        NOT NULL  
);
```

```
CREATE TABLE unit
```

```
(  
    unit_code      CHAR(7)      NOT NULL    PRIMARY KEY,  
    unit_title     VARCHAR(50)  NOT NULL,  
    credit_points  TINYINT     NOT NULL    DEFAULT 15,  
    description    TEXT        NULL  
);
```

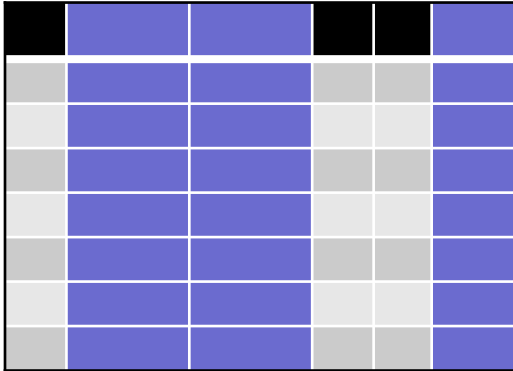
- And then the enrolments table...

```
CREATE TABLE Enrolment  
(  
  enrolment_id    INT           NOT NULL  IDENTITY  PRIMARY KEY,  
  student_id     INT           NOT NULL,  
  unit_code      CHAR(7)       NOT NULL,  
  semester       TINYINT      NOT NULL,  
  year           SMALLINT     NOT NULL,  
  CONSTRAINT stud_fk FOREIGN KEY (student_id) REFERENCES student(student_id),  
  CONSTRAINT unit_fk FOREIGN KEY (unit_code) REFERENCES unit(unit_code)  
);
```

- enrolment_id is primary key, using IDENTITY property to implement an **auto-incrementing integer**
- Foreign key constraints simply specify the FK field, and the PK field that it refers to (inside the name of its home table)

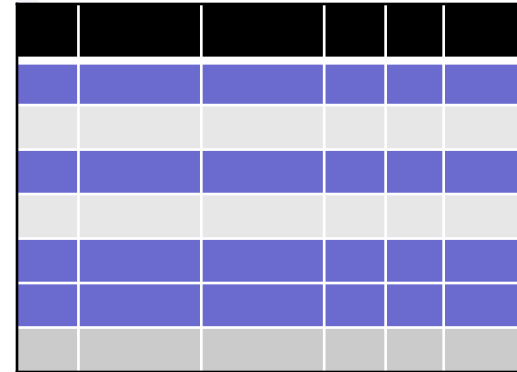
Capabilities of SELECT Statement

Projection



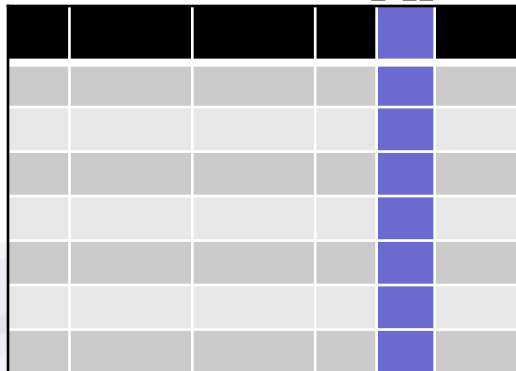
(select certain columns)

Selection



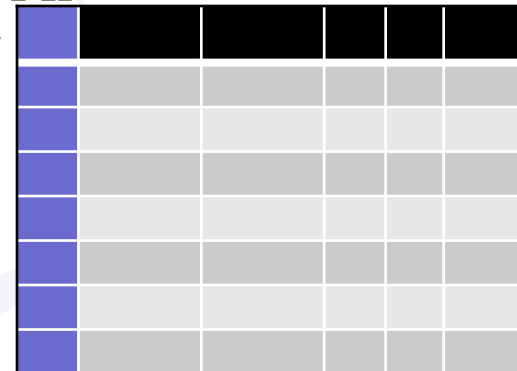
(select certain rows)

FK





PK



Join two tables via keys

Basic SELECT Statement: Syntax

```
SELECT      *|{[DISTINCT] column|expression [alias],...}  
FROM        table  
[WHERE       Conditions];
```

- SELECT clause identifies *which column(s)*
- FROM clause identifies *which table(s)*
- WHERE clause identifies *which row(s)*
 - Where clause is optional
 - Only those rows whose values make the *conditions* true will be returned

Selecting all columns

```
SELECT *  
FROM job;
```

Grid format result:

	job_id	job_title	min_salary	max_salary
1	AC_ACCOUNT	Public Accountant	4200.00	9000.00
2	AC_MGR	Accounting Manager	8200.00	16000.00
3	AD_ASST	Administration Assistant	3000.00	6000.00

Text format result:

job_id	job_title	min_salary	max_salary
-----	-----	-----	-----
AC_ACCOUNT	Public Accountant	4200.00	9000.00
AC_MGR	Accounting Manager	8200.00	16000.00
AD_ASST	Administration Assistant	3000.00	6000.00
AD_PRES	President	20000.00	NULL
AD_VP	Administration Vice President	15000.00	30000.00
IT_PROG	Programmer	4000.00	10000.00
MK_MAN	Marketing Manager	9000.00	15000.00
MK_REP	Marketing Representative	4000.00	9000.00
SA_MAN	Sales Manager	10000.00	20000.00
SA_REP	Sales Representative	6000.00	12000.00
ST_CLERK	Stock Clerk	1000.00	5000.00
ST_MAN	Stock Manager	5500.00	8500.00
(12 row(s) affected)			

Selecting Specific Columns

```
SELECT job_id, max_salary  
FROM   job;
```

<i>job_id</i>	<i>max_salary</i>
AC_ACCOUNT	9000.00
AC_MGR	16000.00
AD_ASST	6000.00
AD_PRES	NULL
AD_VP	30000.00
IT_PROG	10000.00
MK_MAN	15000.00
MK_REP	9000.00
SA_MAN	20000.00
SA_REP	12000.00
ST_CLERK	5000.00
ST_MAN	8500.00

(12 row(s) affected)

Guideline for writing SQL statements

- Goals: correct syntax, readability and easy to edit
 - SQL statements are **not case sensitive**
 - SQL **statements** can be on **one or more lines**
 - **Keywords cannot** be abbreviated or split across lines
 - Keywords are typically entered in **UPPERCASE**
 - **New clauses** are usually placed on separate lines
 - **Indents** are used to enhance **readability**

- Can create expressions with number and date data by using arithmetic operators.
 - Add (+), subtract (-), multiply (*) and divide (/)
 - Other operators exist, but these are the most common
- Operator Precedence:
 - Multiplication and division take priority over addition and subtraction
 - Operators of the same priority are evaluated from left to right
 - Parentheses are used to prioritise evaluation and to clarify statements

```
SELECT job_id, min_salary, max_salary, max_salary*1.05  
FROM job;
```

job_id	min_salary	max_salary	
AC_ACCOUNT	4200.00	9000.00	9450.000000
AC_MGR	8200.00	16000.00	16800.000000
AD_ASST	3000.00	6000.00	6300.000000
AD_PRES	20000.00	NULL	NULL
AD_VP	15000.00	30000.00	31500.000000
IT_PROG	4000.00	10000.00	10500.000000
MK_MAN	9000.00	15000.00	15750.000000
MK_REP	4000.00	9000.00	9450.000000
SA_MAN	10000.00	20000.00	21000.000000
SA_REP	6000.00	12000.00	12600.000000
ST_CLERK	1000.00	5000.00	5250.000000
ST_MAN	5500.00	8500.00	8925.000000
(12 row(s) affected)			

No column name
for "max_salary*1.05"

```
SELECT job_id, min_salary, max_salary, max_salary*1.05
FROM   job;
```

job_id	min_salary	max_salary	
AC_ACCOUNT	4200.00	9000.00	9450.000000
AC_MGR	8200.00	16000.00	16800.000000
AD_ASST	3000.00	6000.00	6300.000000
AD_PRES	20000.00	NULL	NULL
AD_VP	15000.00	30000.00	31500.000000
IT_PROG	4000.00	10000.00	10500.000000
MK_MAN	9000.00	15000.00	15750.000000
MK_REP	4000.00	9000.00	9450.000000
SA_MAN	10000.00	20000.00	21000.000000
SA_REP	6000.00	12000.00	12600.000000
ST_CLERK	1000.00	5000.00	5250.000000
ST_MAN	5500.00	8500.00	8925.000000

(12 row(s) affected)

- A **NULL** is a value that is *unavailable, unassigned, unknown, or inapplicable*
- A null is **not** the same as zero or a blank space
- Arithmetic expressions containing a null value evaluate to null

- A column alias:
 - Renames a column heading for the results of that query
 - Is useful with calculations and other situations where column names may be missing, unhelpful or ambiguous
- Usual form is **AS 'aliasname'** after the column name
 - The AS is optional, but recommended for clarity
 - If the alias contains spaces or special characters, you must enclose it in single quotation marks
 - If the alias is a single word with no special characters, the quote marks can be omitted

Column Aliases

```
SELECT job_title AS Job, max_salary AS 'Pre-raise Maximum',  
       max_salary*1.05 AS 'Post-raise Maximum'  
FROM   job;
```

Job	Pre-raise Maximum	Post-raise Maximum
Public Accountant	9000.00	9450.000000
Accounting Manager	16000.00	16800.000000
Administration Assistant	6000.00	6300.000000
President	NULL	NULL
Administration Vice President	30000.00	31500.000000
Programmer	10000.00	10500.000000
Marketing Manager	15000.00	15750.000000
Marketing Representative	9000.00	9450.000000
Sales Manager	20000.00	21000.000000
Sales Representative	12000.00	12600.000000
Stock Clerk	5000.00	5250.000000
Stock Manager	8500.00	8925.000000

(12 row(s) affected)

- A function named CONCAT() allows you to join multiple text-based columns into a single column of results

```
SELECT CONCAT(job_id, job_title) AS 'Job ID & Title'  
FROM    job;
```

```
Job ID & Title  
-----  
AC_MGRAccounting Manager  
AD_ASSTAdministration Assistant  
AD_VPAdministration Vice President  
MK_MANMarketing Manager  
MK_REPMarketing Representative  
AD_PRESPresident  
IT_PROGProgrammer  
AC_ACCOUNTPublic Accountant  
SA_MANSales Manager  
SA_REPSales Representative  
ST_CLERKStock Clerk  
ST_MANStock Manager  
  
(12 row(s) affected)
```

- You can add other text to this *in single quotes as needed*

```
SELECT CONCAT(job_id, ' (' , job_title, ')') AS 'Job ID & Title'  
FROM    job;
```

Job ID & Title

AC_MGR (Accounting Manager)
AD_ASST (Administration Assistant)
AD_VP (Administration Vice President)
MK_MAN (Marketing Manager)
MK_REP (Marketing Representative)
AD_PRES (President)
IT_PROG (Programmer)
AC_ACCOUNT (Public Accountant)
SA_MAN (Sales Manager)
SA_REP (Sales Representative)
ST_CLERK (Stock Clerk)
ST_MAN (Stock Manager)

(12 row(s) affected)

- A common use of concatenation is to produce full names...

```
SELECT CONCAT(first_name, ' ', last_name) AS 'full_name'  
FROM   employee;
```

<i>Full_name</i>

<i>Steven King</i>
<i>Neena Kochhar</i>
<i>Lex De Haan</i>
<i>Alexander Hunold</i>
<i>Bruce Ernst</i>
<i>...</i>
<i>(20 row(s) affected)</i>

- In earlier versions of SQL Server, “+” was used to concatenate. This is still supported (only in SQL Server)

```
SELECT first_name + ' ' + last_name AS 'full_name'  
FROM   employee;
```

Duplicate Rows & the DISTINCT Keyword

- *By default, a query will display all rows, including rows which contain the same values (i.e. duplicates)*

```
SELECT job_id  
FROM   employee;
```

```
job_id  
-----  
AD PRES  
AD VP  
AD VP  
IT PROG  
IT PROG  
IT PROG  
ST MAN  
ST CLERK  
ST CLERK  
...  
(20 row(s) affected)
```

Duplicate Rows & the DISTINCT Keyword

- You can eliminate duplicate rows by using the **DISTINCT** keyword in the **SELECT** clause.*

```
SELECT DISTINCT job_id  
FROM   employee;
```

job_id


AC_ACCOUNT
AC_MGR
AD_ASST
AD PRES
AD_VP
IT_PROG
MK_MAN
MK_REP
SA_MAN
SA_REP
ST_CLERK
ST_MAN

(12 row(s) affected)

- That covers the basics of the SELECT statement
 - Selecting all columns with *
 - Specifying which columns to select by column name
 - Arithmetic operators
 - Column aliases
 - Concatenation
 - Using the DISTINCT keyword to eliminate duplicates
- These can all be combined in order to select something very specific from a table
- In coming weeks, we will learn to specify criteria with the WHERE clause, and connect columns from different tables using joins

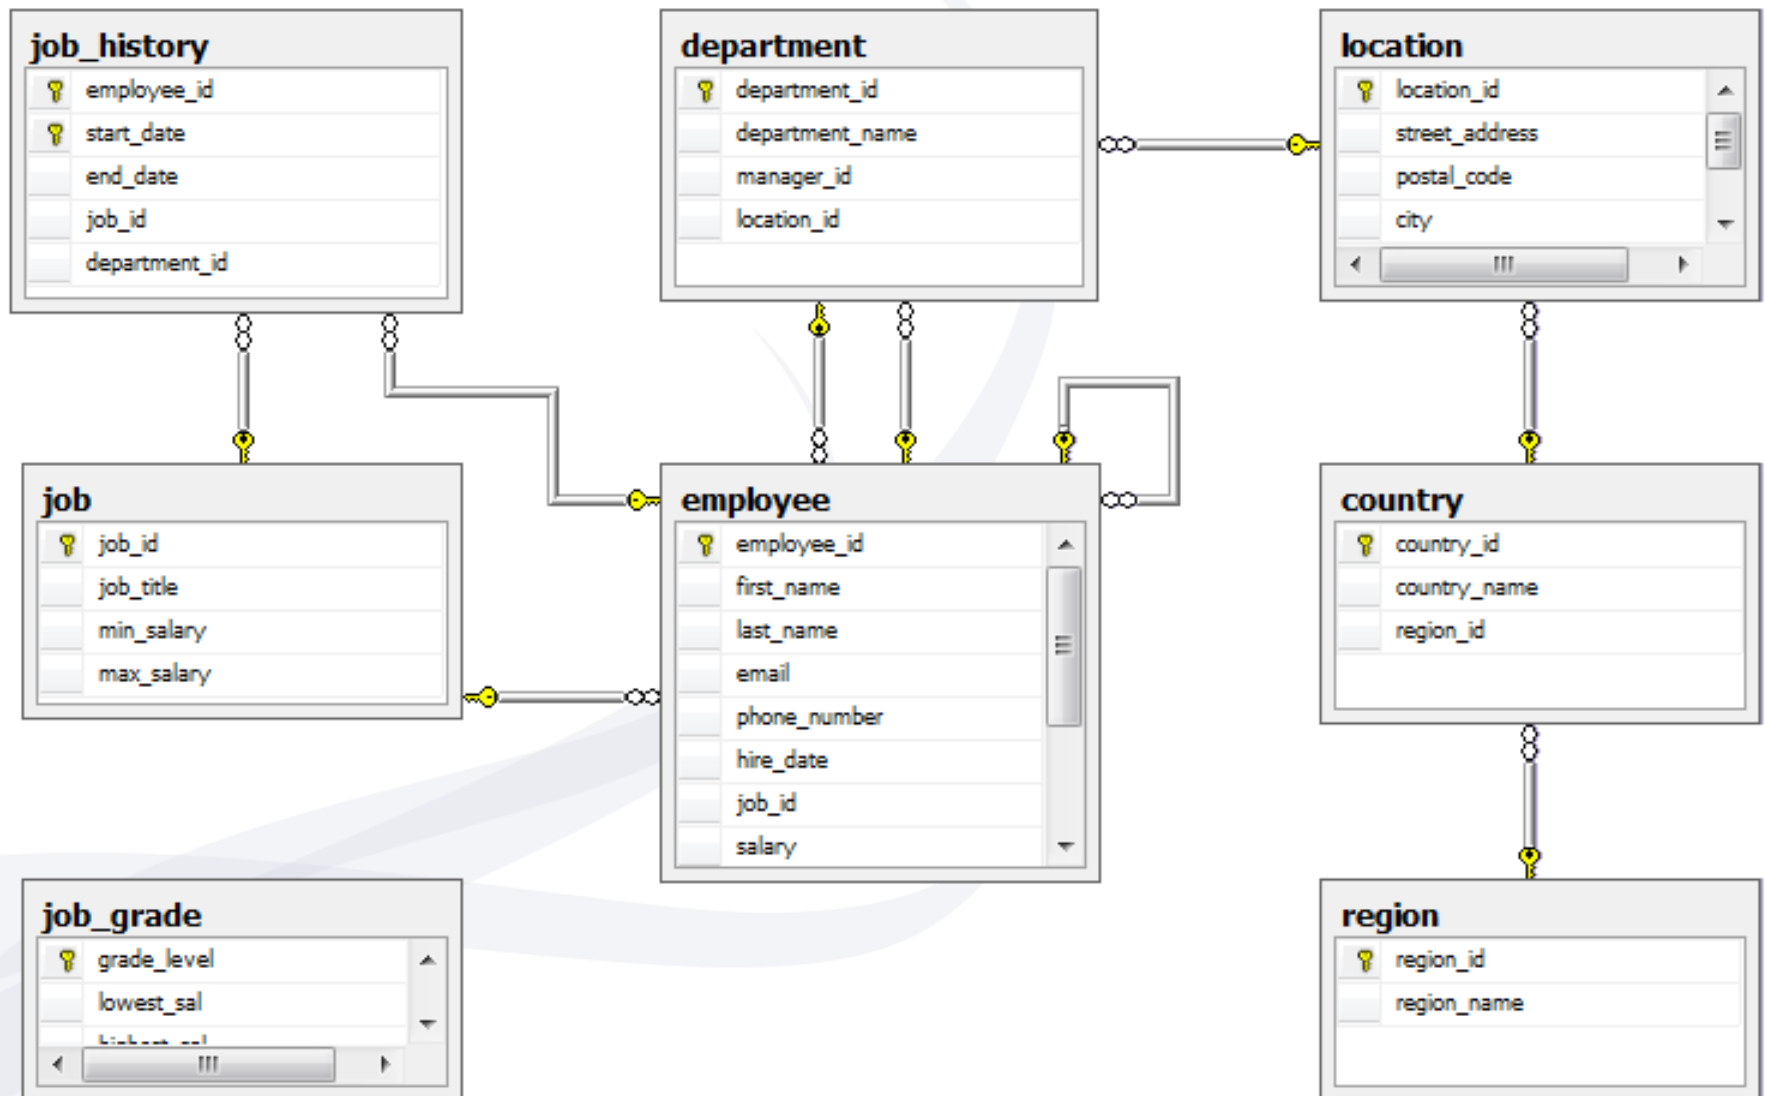
- Microsoft SQL Server includes multiple components and services that makes it a comprehensive enterprise platform
- Key Components:
 - Database Engine
 - Analysis Services
 - Integration Services
 - Replication Services
 - Reporting Services
 - ***Service Broker***
 - ***Native HTTP Support***
 - ***.NET Common Language Runtime (SQL CLR)***
 - ***Notification Services***
 - ***Full-Text Search, and more!***
- Newer versions of SQL Server exist, but the implementation of SQL in the newer versions is almost identical
 - This unit focuses on the SQL, not the other server features

- SSMS is an integrated database management and development environment for SQL Server
- Developers can use it to
 - Create databases, tables and other objects
 - Change table structures and constraints
 - Create, execute and save scripts/queries
 - Manage databases, services and other components
 - Create database diagrams



Company ER Diagram in SSMS

- SSMS can generate database diagrams such as this...



SQL Server Management Studio (SSMS)

The screenshot displays the Microsoft SQL Server Management Studio (SSMS) interface. The title bar indicates the connection to 'company.sql - GREGOFFICE.company (GREGOFFICE\Greg (52))'. The menu bar includes File, Edit, View, Query, Project, Tools, Window, and Help. The toolbar contains icons for various actions like opening files, saving, and executing queries. The Object Explorer on the left shows the server structure for 'GREGOFFICE (SQL Server 12.0.2269 - GR...)', with folders for Databases, System Databases, company, Database Diagrams, Tables, Views, Synonyms, Programmability, Service Broker, Storage, Security, SalesSystem, and Server Objects. The Working Area in the center shows a SQL script with the following content:

```
DROP DATABASE company;  
END  
GO  
  
-- Now that we are sure the database does not exist, we create it.  
  
PRINT 'Creating database.';  
CREATE DATABASE company;  
GO  
  
-- Now that an empty database has been created, we will make it the active one.  
-- The table creation statements that follow will therefore be executed on the newly created database.  
  
USE company;  
GO
```

The Query Results pane at the bottom shows the output of the executed query:

```
Database exists - dropping.  
Creating database.  
Creating region table...  
Creating country table...  
Creating location table...  
Creating job table...  
Creating department table...  
Creating employee table...  
Creating job_history table...  
Creating job_grade table...  
Populating region table...  
  
(4 row(s) affected)  
Populating country table...  
  
(4 row(s) affected)  
Populating location table...
```

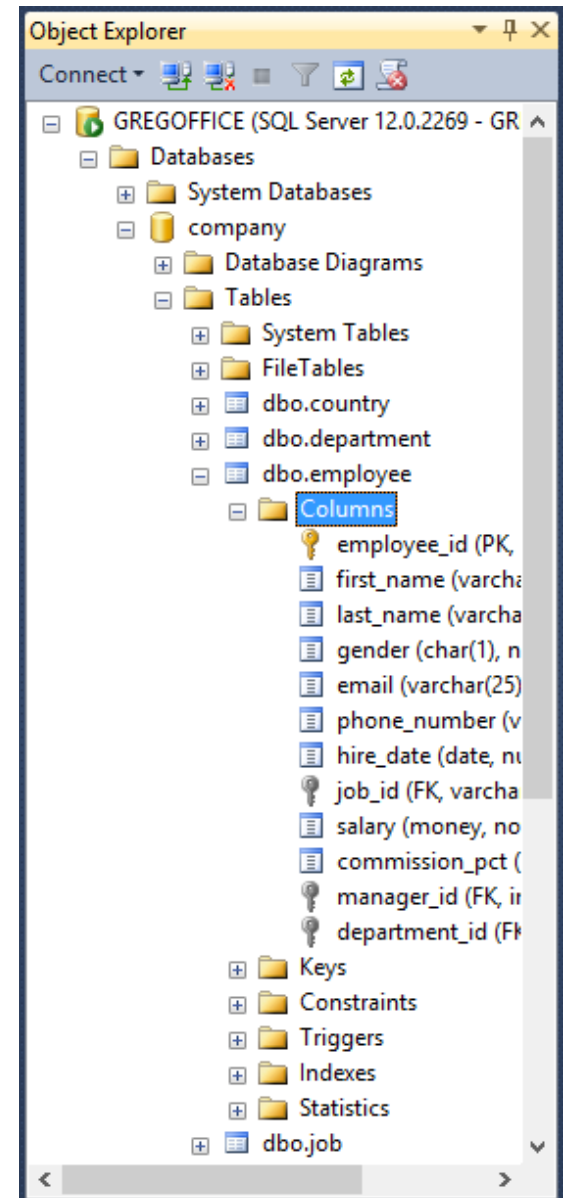
The status bar at the bottom indicates 'Query executed successfully.' and provides details about the connection and execution time.

Object Explorer

Working Area

Query Results

- The left menu is the **Object Explorer**
- **Browse**, access and manage all objects (tables, databases, etc) on the server
- **Refresh button** may be needed before new items show up
- **Right click** on objects for menu of useful/common commands
 - Right click on a table (or view) to select or edit its contents
 - Right click on column to edit data type, length, null, etc



- Many of the toolbar buttons are very useful, but here are a few you will probably use the most frequently
 - Some toolbars will only appear while working on a query
 - Hover over other buttons to see tooltip about their purpose

**Save active tab
in the work area**

**Open new query
tab in work area**



Select active database

**Run (selected part
of or whole) query**

**Show results
as table or text**

- **Note:** If you have some text selected in your query, *only the selected text* will be executed when you press the execute button

- In ECU labs:
 - Find it by searching for it in the start menu (type “SQL”) and connect with the default settings (just press Connect button)
 - *Follow the instructions in Tutorial 7 to get started this week*
- At home:
 - *Follow the installation guide in Canvas*
 - Download and install SQL Server 2014 Express Edition
 - This is a free version of SQL Server 2014
 - Make sure you install Management Studio
 - Once installed, launch Management Studio and connect with the default settings
- Once set up, you're ready to play with some databases!

Summary

- Introduction to SQL
- Basics of the SELECT statement
- Introduction to SQL Server and SQL Server Management Studio (SSMS)
- From this week onwards the workshop sessions are labs, in which we work in SSMS

Questions?